

Agenda: Azure Storage Service

- Introduction to Storage Service
- Creating a Storage Account
- Working with Blob Storage
- Azure Storage Explorer
- Transfer Data using AzCopy
- Azure blob storage Lifecycle Management
- Storage account Security
- Working with Table Storage
- Azure SMB File Storage

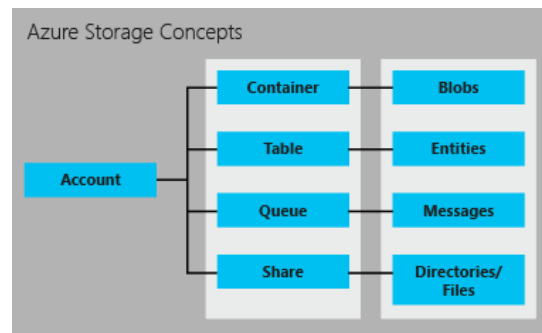
Introduction to Azure Storage Service

Cloud computing enables new scenarios for applications requiring **scalable, durable and highly available** storage for their data – which is exactly why Microsoft developed **Azure Storage Service**.

- Azure Storage is a **PaaS service** that you can use to store both **unstructured** and **partially structured** data.
- **Azure Storage is massively scalable and elastic:** It can store and process **hundreds of terabytes of data** to support the big data scenarios required by scientific, financial analysis, and media applications. Or you can store the **small amounts of data** required for a small business website.
- By default, you can create up to **100 storage accounts** in a single Azure subscription. Each standard storage account can contain up to **500 TB** of combined blob, queue, table and file data.
- As the demands on your storage application grow, Azure Storage **automatically allocates** the appropriate resources to meet them. **We are charged only for what we use.**
- SDKs for Azure Storage in a variety of languages – **.NET, Java, Node.js, Python, PHP, Ruby, Go**, and others – as well as **REST APIs**. It also supports scripting in Azure PowerShell or Azure CLI.

It offers **four types of storage services**, depending on the type of data that they are designed to store:

1. **Blob Storage** stores file data. A blob can be any type of **text or binary data**, such as a document, media file, database backup, VHD files or application installer. Blob Storage is sometimes referred to as **Object storage**.
2. **File Storage** Similar to blobs, these provide storage for unstructured files, but they offer support for file sharing in the same manner as traditional on-premises Windows file shares.
3. **Table Storage** stores partially structured datasets. Table storage is a **NoSQL** key-attribute data store, which allows for rapid development and fast access to large quantities of data.
4. **Queue Storage** provides **reliable messaging** for workflow processing and for communication between components of cloud services.



Azure Storage Account

An Azure storage account is a **secure account** that gives you access to services in Azure Storage..Your storage account provides the unique namespace for your storage resources. There are two types of storage accounts:

1. A **standard storage** account includes Blob, Table, Queue, and File storage. Standard use **HDD** Drives
2. A **premium storage** account is ideally supposed to be used for Azure Virtual Machine disks. Premium use **SSD** Drives

*Note:*Combining data services(Blob, Table, Queue, and File storage) into a storage account lets you manage them as a group. The settings you specify when for the account, are applied to everything in the account. Deleting the storage account deletes all of the data stored inside it.

Creating Storage Account

The tools available for creating storage account are Portal,CLI,Powershell,Management client libraries.

The Azure CLI and Azure PowerShell let you write scripts, while the management libraries allow you to incorporate the creation into a client app.

Lab1: Create Storage Account Using Portal

1. Azure Portal → Search → **Storage accounts** → **+Create**

Basic tab:

- a) Storage account name = sandeepdemostorage (must be all lowercase)
- b) Performance: Standard / Premium
- c) Redundancy = LRS

Advanced Tab:

- a) Allow enabling anonymous access on individual containers = Check

Other Tabs:

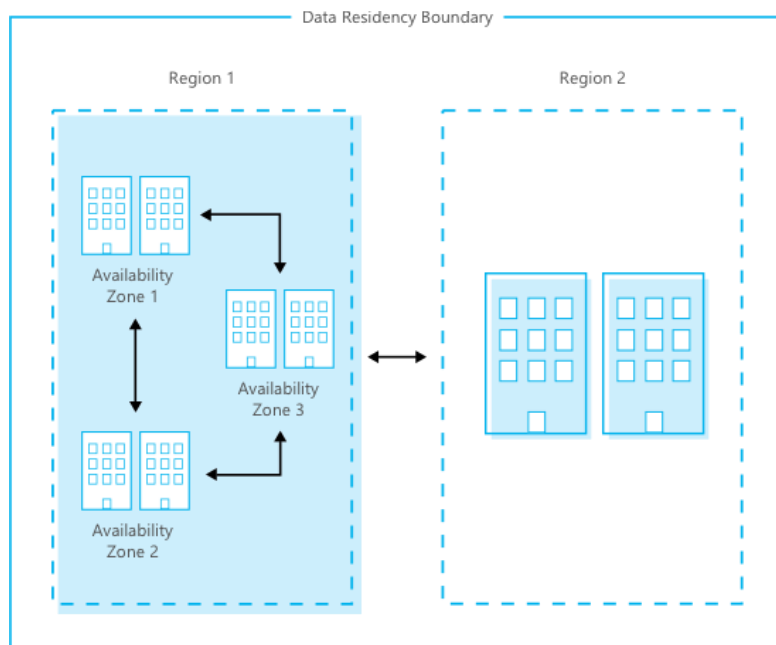
- a) Read and pass on – Leave them as default
2. Review → Create

About Account Kind:

- **General-purpose v2 accounts:** Basic storage account type for blobs, files, queues, and tables. Recommended for most scenarios using Azure Storage.
- **Block blob storage accounts:** Blob-only storage accounts with premium performance characteristics. Recommended for scenarios with high transactions rates, using smaller objects, or requiring consistently low storage latency.
- **FileStorage storage accounts:** Files-only storage accounts with premium performance characteristics. Recommended for enterprise or high performance scale applications.

Storage account type	Supported services	Supported performance tiers	Replication options
General-purpose V2	Blob, File, Queue, Table, and Disk (VHD)	Standard	LRS, GRS, RA-GRS, ZRS, ZGRS RA-ZGRS
Block blob storage	Blob (block blobs* and append blobs** only)	Premium* Standard**	LRS / ZRS
FileStorage	Files only	Premium	LRS / ZRS
Page Blobs	Disk (VHD)	Premium	LRS/ ZRS

Azure Regions and availability Zones :



Regions

- A region is a set of datacenters deployed within a latency-defined perimeter and connected through a dedicated regional low-latency network.

Availability Zones

- Availability Zones are unique physical locations within an Azure region. Each zone is made up of one or more datacenters equipped with independent power, cooling, and networking.

About Redundancy

Locally redundant storage (LRS):

- **Replicates 3 times within a single data center in a single region where Storage Account is created.**
- Locally redundant storage (LRS) provides at least 99.999999999% (11 nines) durability of objects over a given year.
- The replicas are spread across UD and FDs within one storage scale unit (A storage scale unit is a collection of racks of storage nodes.)
- A request returns successfully only once it has been written to all three replicas.
- This architecture ensures your data is available if a hardware failure affects a single rack or when nodes are upgraded during a service upgrade.
- LRS is less expensive than GRS and also offers higher throughput.
- For Premium Storage accounts - This is the only option available.

Zone-redundant storage (ZRS)

- Replicates your data across three (3) storage clusters in a single region. Each storage cluster is physically separated from the others and resides in its own availability zone. Each availability zone, and the ZRS cluster within it, is autonomous, with separate utilities and networking capabilities.
- ZRS is not yet available in all regions.
- Once you have created your storage account and selected ZRS, you **cannot convert** it to use to any other type of replication, or vice versa.
- Consider ZRS for scenarios that require strong consistency, strong durability, and high availability even if an outage or natural disaster renders a zonal data center unavailable.
- **What happens when a zone becomes unavailable?** Your data is still accessible for both read and write operations even if a zone becomes unavailable. Microsoft recommends that you continue to follow practices for **transient fault handling**. These practices include implementing retry policies with exponential back-off.

Geo-redundant storage (GRS)

- GRS maintains **6 copies** of your data. 3 replicas in **primary region** and also replicates your data 3 additional times to a **secondary region** that is hundreds of miles away from the primary region.
- 99.99999999999999 (16 9s) Availability is supported when calculated around a year.

- Data is durable even in the case of a complete regional outage or a disaster in which the primary region is not recoverable.
- If failure occurs in the primary region, Azure Storage **automatically failover** to the secondary region.
- An update is first committed to the primary region, where it is replicated three times. Then the update is replicated to the secondary region, where it is also replicated three times.
- Requests to write data are replicated **asynchronously** to the secondary region. It is important to note that opting for GRS does not impact latency of requests made against the primary region.
- The secondary region is **automatically determined** based on the primary region, and cannot be changed.

Read-access geo-redundant storage (RA-GRS)

- As with GRS, your data replicates asynchronously across two regions and synchronously within each region, yielding six copies of a storage account.
- This is default option when we create a storage account.
- In the event that data becomes unavailable in the primary region, your application can **read data** from the secondary region.
- If your primary endpoint for the Blob service is **myaccount.blob.core.windows.net**, then your secondary endpoint is **myaccount-secondary.blob.core.windows.net**. The **access keys** for your storage account are the **same** for both the primary and secondary endpoints.

More Details: <https://azure.microsoft.com/en-in/documentation/articles/storage-redundancy/>

About Access tier: Cool / Hot

- Access tier: **Hot**, if objects will be **more** frequently accessed. This allows you to store data at a **lower access cost. Higher Storage Cost.**
- Access tier: **Cool**, if objects will be **less** frequently accessed. This allows you to store data at a **lower data storage cost. Higher RW Cost. Min 30 days.**
- Access tier: **Cold**, if objects will be **less** frequently accessed. This allows you to store data at a **lower data storage cost. Higher RW Cost. Min 90 days.**
- Access tier: **Archive**. The archive tier is optimized for data that can tolerate **several hours of retrieval latency** and will remain in the Archive tier for at least 180 days. The archive tier is the most cost-effective option for storing data, but accessing that data is more expensive than accessing data in the hot or cool tiers. It is available at level of an individual blob only, not at the storage account level. Only block blobs and append blobs can be archived. Min 180 days.

Rehydrate an archived blob to an online tier:

- To read data in archive storage, you must first change the tier of the blob to hot or cool. This process is known as rehydration and can take hours to complete.

- There are currently two rehydrate priorities, High and Standard, which can be set via the optional **x-ms-rehydrate-priority** property on a **Set Blob Tier** or **Copy Blob** operation.
 - **Standard priority:** The rehydration request will be processed in the order it was received and **may take up to 15 hours**.
 - **High priority:** The rehydration request will be prioritized over Standard requests and may finish **in under 1 hour**.

Secure transfer required = **Disabled** / Enabled

If Enabled only HTTPS requests will be accepted.

This option doesn't work with Custom Domain Names for Storage account.

Pricing and Billing of Storage Account

All storage accounts use a pricing model for blob storage based on the tier of each blob.

When using a storage account, the following billing considerations apply:

- **Storage costs:** In addition to, the amount of data stored, the cost of storing data varies depending on the storage tier. The per-gigabyte cost decreases as the tier gets cooler.
- **Data access costs:** Data access charges increase as the tier gets cooler. For data in the cool and archive storage tier, you are charged a per-gigabyte data access charge for reads.
- **Transaction costs:** There is a per-transaction charge for all tiers that increases as the tier gets cooler.
- **Geo-Replication data transfer costs:** This charge only applies to accounts with geo-replication configured, including GRS and RA-GRS. Geo-replication data transfer incurs a per-gigabyte charge.
- **Outbound data transfer costs:** Outbound data transfers (data that is transferred out of an Azure region) incur billing for bandwidth usage on a per-gigabyte basis, consistent with general-purpose storage accounts.
- **Changing the storage tier:** Changing the account storage tier from cool to hot incurs a charge equal to reading all the data existing in the storage account. However, changing the account storage tier from hot to cool incurs a charge equal to writing all the data into the cool tier (GPv2 accounts only).

Working with Blob Storage

- **Blobs** are binary large objects. The Blob service stores text and binary data.
- Blob storage is also referred to as **object storage**.
- Every blob is organized into a **container**. Containers also provide a useful way to assign security policies to groups of objects. A storage account can contain any number of containers, and a container can contain any number of blobs, up to the **500 TB capacity** limit of the storage account.
- **Creating BLOB Hierarchies:** The blob service in Azure Storage is based on a **flat storage scheme**. This means that creating a container one level below the **root is the only true level of container**. However, you can specify a delimiter as part of the blob name to create your own **virtual hierarchy**. For example, you could create a blob

named **/January/Reports.txt** and **/February/Reports.txt**, and filter based on **/January** or **/February** in most tools that support Azure Storage. Most third-party storage tools allow you to create folders within a container, but they are actually being clever with the name of the blob itself.

Types of blobs:

1. **Block blobs** are optimized for streaming (**sequential access**) and for uploads and downloads, and are a good choice for storing documents, media files, backups etc. Azure divides data into smaller blocks of up to 100 megabytes (MB) in size, which subsequently upload or download in parallel. Individual block blobs (file) can be up to 100 GB in size. One blob can have max of 50,000 blocks. (50,000 blocks * 4 MiB each = 200,000 MiB = 200 GB)
2. **Append blobs**: Append blobs are similar to block blobs, but are optimized for append operations. This works best with **logging and auditing** activities. Updating or deleting of existing blocks is not supported.
3. **Page blobs** are optimized for **random read/write** operations and provide the ability to write to a range of bytes in a blob. Blobs are accessed as pages, each of which is up to **512 bytes** in size. Each Page blob can be up to **8TB** each. Is best suited for **virtual machine disks**.

Blobs are addressable using the following URL format:

http(s)://<storage account name>.blob.core.windows.net/<container>/<blob name>

Lab2: Create container, Upload/download blobs, work with soft delete, snapshot, public and private blobs.

1. Ensure that Public Access is Enabled

Storage Account → Configuration → Enable **Allow enabling anonymous access on individual containers**

2. Create container with Anonymous blob access

Storage Account → Containers → +Container

- Name: **public-images**
- Anonymous access level: Blob(anonymous read access for blobs only)
- Create

3. Upload File

Select container, public-images → Upload → Browse for files (select any file/image) → Upload

4. Open File URL in browser

Select the blob → Copy URL → Open URL in browser.

Lab3: Create container, Upload/download blobs, work with soft delete, snapshot, public and private blobs.

1. Create container with private blob access

Storage Account → Containers → +Container

- Name: **private-images**

- Anonymous access level: private(no anonymous access)
- Create

2. Upload File

Select container, private-images → Upload → Browse for files (select any file/image) → Upload

3. Open File URL in browser

Select the blob → Copy URL → Open URL in browser → Note the error

4. Get the Access Token

Select the blob → Generate SAS Tab → Click **Generate SAS token and URL** → Scroll down and Copy **Blob SAS URL**

5. Open File URL in browser

Open URL in browser → You can view the file

Soft Delete (Blob Service → Data Protection → Blob soft delete)

- Enables blobs and checkpoints to be recovered if deleted.
- It utilizes special soft delete snapshot.
- Configured at storage account level for specific retention duration.

Immutable blob:

- It enables blob to be configured as read only so that it can not be altered or deleted.
- It is configured at container level.
- Steps to configure:
 - a. Select Container → Access Policy → Immutable blob storage → + Add policy
 - b. Now, you wont be able able to delete any blob.

Update retention period to *

1 days

☐ Enable version-level immutability ⓘ

Allow protected append writes to ⓘ

☐ None

☐ Append blobs

☒ Block and append blobs

Blob snapshot:

- Enables point in time copy of blob content to be taken.
- They are associated with parent blob.

Types of blobs:

1. **Block blobs** are optimized for streaming (**sequential access**) and for uploads and downloads, and are a good choice for storing documents, media files, backups etc. Azure divides data into smaller blocks of up to 100

megabytes (MB) in size, which subsequently upload or download in parallel. Individual block blobs (file) can be up to 100 GB in size. One blob can have max of 50,000 blocks.

2. **Append blobs:** Append blobs are similar to block blobs, but are optimized for append operations. This works best with **logging and auditing** activities. Updating or deleting of existing blocks is not supported.
3. **Page blobs** are optimized for **random read/write** operations and provide the ability to write to a range of bytes in a blob. Blobs are accessed as pages, each of which is up to **512 bytes** in size. Each Page blob can be up to **8TB** each. Is best suited for **virtual machine disks**.

Azure Storage Explorer

Microsoft Azure Storage Explorer is a standalone app from Microsoft that allows you to easily work with Azure Storage data.

Some of the benefits of Azure Storage Explorer are:

- a) Access multiple accounts and subscriptions across Azure
- b) Create, delete, view, and edit storage resources
- c) View and edit Blob, Queue, Table, File, Cosmos DB storage and Data Lake Storage.
- d) Obtain shared access signature (SAS) keys
- e) Available for Windows, Mac, and Linux

To download storage explorer:

<https://download.microsoft.com/download/A/E/3/AE32C485-B62B-4437-92F7-8B6B2C48CB40/StorageExplorer-windows-x64.exe>

Transfer data with the AzCopy

- AzCopy is a command-line utility designed for copying data to/from Microsoft Azure Blob, File, and Table storage, using simple commands designed for optimal performance.
- You can copy data between a file system and a storage account, or between storage accounts.
- *The AzCopy Tool supports a maximum file size of 1Tb and will automatically split into multiple files if the data file exceeds 200Gb.*

There are two versions of AzCopy that you can download.

1. **AzCopy on Windows** is built with .NET Framework, and offers Windows style command-line options.
2. **AzCopy on Linux** is built with .NET Core Framework which targets Linux platforms offering POSIX style command-line options.

The basic syntax for AzCopy commands is:

```
AzCopy /Source:<source> /Dest:<destination> [Options]
```

Lab4: Using AzCopy

Step1: Download AZCopy

1. Download Zip file from <https://learn.microsoft.com/en-us/azure/storage/common/storage-use-azcopy-v10#download-azcopy>
2. Extract the Zip file in folder D:\azcopy

Step2: Login to Azure Storage:


- a) Azure Portal → Azure Active Directory → Copy Tenant ID
- b) D:\azcopy>azcopy login --tenant-id <Paste Tenant ID here>

Note: The login user should have **Storage Blob Data Reader / Storage Blob Data Owner / Storage Blob Data Contributor** Role

Step3: Assign yourself Storage Blob Data Contributor Role

Storage Account → Access Control (IAM) → +Add → Add role assignment →

- a) **Role Tab:** Search and Select **Storage Blob Data Contributor** → Next
- b) **Members Tab:** + Select members → search and select **your login identity** → Select
- c) Review + assign → Assign



Step4: Copy single file (d:\demo.txt) from local machine to blob storage:

```
D:\azcopy>azcopy copy "d:\demo.txt" "https://<Storage Account Name>.blob.core.windows.net/public-images/demo.txt"
```

Step4: Copy complete folder (d:\images) from Local to blob storage:

```
D:\azcopy>azcopy copy "d:\images" "https://<Storage Account Name>.blob.core.windows.net/public-images" --recursive=true
```

Step5: Copy from one storage account to another blob storage:

```
D:\azcopy>azcopy copy "https://<Source Storage Account Name>.blob.core.windows.net/source" "https://<Target Storage Account Name>.blob.core.windows.net/target?<SAS Token>" --recursive=true
```

Refer for more information:

<https://docs.microsoft.com/en-us/azure/storage/common/storage-use-azcopy-v10>

Azure blob storage Lifecycle Management

Azure Blob storage lifecycle management offers policy to transition your data to the appropriate access tiers (Hot,Cool,Archive)or expire at the end of the data's lifecycle.

By adjusting storage tiers in respect to the age of data, you can design the least expensive storage options for your needs. You can create a policy that deletes old snapshots based on snapshot age.

The lifecycle management policy lets you:

- Transition blobs to a cooler storage tier (hot to cool, hot to archive, or cool to archive) to optimize for performance and cost
- Delete blobs at the end of their lifecycles
- Define rules to be run once per day at the storage account level
- Apply rules to containers or a subset of blobs (using prefixes as filters)

Example:

- Change blob to Cool tier 30 days after last modification
- Change blob to Archive tier 90 days after last modification
- Delete blob 2,555 days (seven years) after last modification
- Delete blob snapshots 90 days after snapshot creation

Lab5: Create Lifecycle management policy for blobs

Storage Account → Lifecycle Management

Add a rule

[Action set](#) [Filter set](#) [Review + add](#)

Each rule definition includes an action set and a filter set. The action set applies the tier or delete actions to the filtered set of objects. The filter set limits rule actions to a certain set of objects within a container or objects names.

Rule name * ✓

Blobs

☒ Move blob to cool storage
Days after last modification ✓☒ Move blob to archive storage
Days after last modification ✓☒ Delete blob
Days after last modification ✓

Snapshots

☒ Delete snapshot
Days after blob is created ✓

Code View:

```
{
  "rules": [
    {
      "name": "ruleFoo",
      "enabled": true,
      "type": "Lifecycle",
      "definition": {
        "filters": {
          "blobTypes": [ "blockBlob" ],
          "prefixMatch": [ "container1/foo" ]
        },
        "actions": {
          "baseBlob": {
            "tierToCool": { "daysAfterModificationGreaterThan": 30 },
            "tierToArchive": { "daysAfterModificationGreaterThan": 90 },
            "delete": { "daysAfterModificationGreaterThan": 2555 }
          },
          "snapshot": {
            "delete": { "daysAfterCreationGreaterThan": 90 }
          }
        }
      }
    }
  ]
}
```

```
]
}
```

Storage Account Security

Settings → Access Keys:

- Azure creates two of these keys (primary and secondary) for each storage account you create. The keys give access to everything in the account.
- The client can embed the Access key in the HTTP Authorization header of every request, and the Storage account validates the key.
- It is recommended to use these keys only with trusted in-house applications that you control completely.

Storage Account → Configuration → Ensure Allow storage account key access = **Enabled**




Role-based Access control

- Azure Storage supports Azure Active Directory and role-based access control (RBAC) for both resource management and data operations.
- To security principals, you can assign RBAC roles that are scoped to the storage account.

Lab 6: Give RBAC to user for storage account

- a) Storage Account → Configuration → Ensure **Default to Azure Active Directory authorization in the Azure portal = Disabled**
- b) Storage account → Containers → **private-images** → click on [\(Switch to Azure AD User Account\)](#) and note that you are **NOT** able to access data in the container



Note: If you have completed **Lab4: Using AzCopy** you will not get any error and following steps can be ignored.

- c) Storage Account → Access Control (IAM) → +Add → **Add role assignment** →
 - a. Role Tab → Search and Select **Storage Blob Data Contributor** → Next
 - b. Members Tab: + **Select members** → search and select your login Identity → Select
 - c. Review + assign → Assign
- d) Storage account → Containers → **private-images** → click on [\(Switch to Azure AD User Account\)](#) and note that you are able to access data in the container

Shared Access Policy and Shared Access Signature (SAS) Token

A **shared access signature (SAS)** is a URI that grants **restricted access rights** to Azure Storage resources. You can provide a shared access signature to clients who should not be trusted with your storage account key but whom you wish to delegate access to certain storage account resources. By distributing a shared access signature URI to these clients, you grant them access to a resource for a specified **period of time**.

Azure Storage supports two kinds of shared access signatures:

1. An **Account SAS** delegates access to resources in **one or more** of the storage services. You can also delegate access to read, write, and delete operations on blob containers, tables, queues, and file shares that are not permitted with a service SAS.
2. The **Service SAS** delegates access to a resource in **just one** of the storage services: Blob, Queue, Table, or File service.

Note that stored access policies are currently not supported for an account-level SAS.

Lab7: Creating an Account SAS (for many operations)

1. Storage Account → Shared access signature
2. Provide the options as required → **Generate SAS and connection string**

Allowed services: ☒ Blob ☒ File ☐ Queue ☐ Table

Allowed resource types: ☐ Service ☒ Container ☒ Object

Allowed permissions: ☒ Read ☐ Write ☐ Delete ☐ List ☐ Add ☐ Create ☐ Update ☐ Process ☐ Immutable storage ☐ Permanent delete

Blob versioning permissions: ☐ Enables deletion of versions

Allowed blob index permissions: ☐ Read/Write ☐ Filter

Start and expiry date/time: Start: 09/11/2023 1:10:14 PM, End: 09/11/2023 9:10:14 PM

Allowed IP addresses: For example, 168.1.5.65 or 168.1.5.65-168.1.5.70

Allowed protocols: ☒ HTTPS only ☐ HTTPS and HTTP

Preferred routing tier: ☒ Basic (default) ☐ Microsoft network routing ☐ Internet routing

Signing key: key1

Generate SAS and connection string

3. Copy the **SAS token** and share it with the client.

<https://dssdemostorage.blob.core.windows.net/private-images/demo.gif?sv=2017-11-09&ss=b&srt=co&sp=r&se=2018-08-13T10:05:57Z&st=2018-08-11T02:05:57Z&spr=https&sig=hFYQeZ2fvj52%2BQI0kg%2BbPmErr8J%2F5hKrGkAnF7Q7u%2F4%3D>

Stored Access Policies

A Shared Access Signature can take one of two forms:

- a) **An ad hoc SAS**. When you create an ad hoc SAS, the start time, expiration time, and permissions for the SAS are all specified on the SAS URI (or implied in the case where the start time is omitted). This type of SAS can be created on a container, blob, table, or queue.

- b) **An SAS with a Stored Access Policy.** A stored access policy is defined on a resource container—a blob container, table, or queue—and can be used to manage constraints for one or more Shared Access Signatures. When you associate an SAS with a stored access policy, the SAS inherits the constraints—the start time, expiration time, and permissions—defined for the stored access policy.

Note: Stored access policies give you the option to revoke permissions without having to regenerate the storage account keys. Set the expiration on these to be a very long time (or infinite), and make sure that it is regularly updated to move it further into the future.

Lab8: Create Ad-hoc **Service SAS** using Portal (Same as Lab2)

1. Storage Account → Containers → Select private-images → <Select Blob Item> → **Generate SAS Tab** → provide required options → Click on **Generate SAS token and URL**
2. Copy the Blob SAS URL and open in browser

Note: SAS Token without policy cannot be revoked.

Lab 9: Create Stored Access Policy using Portal

1. Storage Accounts → Containers → Select the container private-images → **Access Policy** → + Add Policy (Stored access policies)
2. Enter details as below → OK
3. **Save**

Add policy

Identifier *	Permissions
DemoPolicy ✓	Read ▼
Start time	Expiry time
MM/DD/Y... h:mm:ss AM/PM	MM/DD/Y... h:mm:ss AM/PM
(UTC+05:30) Chennai, Kolkata... ▼	(UTC+05:30) Chennai, Kolkata... ▼

OK Cancel

Create SAS Token using Stored Access Policy

1. Storage Accounts → Containers → Select the container private-images → Select a file/blob
2. Generate SAS tab
 - a. Stored access policy = DemoPolicy

Overview Versions Snapshots Edit **Generate SAS**

A shared access signature (SAS) is a URI that grants restricted access to an Azure Storage blob. Use it when you want to share access to your storage account key. [Learn more about creating an account SAS](#)

Signing method
☒ Account key ☐ User delegation key

Signing key
 Key 1

Stored access policy
 DemoPolicy

Permissions
 Read

Start and expiry date/time

Start
 09/12/2023 3:59:06 PM
 (UTC+05:30) Chennai, Kolkata, Mumbai, New Delhi

Expiry
 09/12/2023 11:59:06 PM
 (UTC+05:30) Chennai, Kolkata, Mumbai, New Delhi

Allowed IP addresses
 for example, 168.1.5.65 or 168.1.5.65-168.1.5.65

Allowed protocols
☒ HTTPS only ☐ HTTPS and HTTP

Generate SAS token and URL

- b. Scroll a bit and copy Blob SAS Url → Paste in Notepad and observe Stored Access Policy is used in URL
- c. Paste the same in Browser to download the file.

Expire the Token

3. Storage Accounts → Containers → Select the container private-images → Access Policy → DemoPolicy → click on ... → Edit
4. Change the Expiry time = <Yesterday date> → OK
5. Save

Edit policy

Identifier * DemoPolicy Permissions Read

Start time MM/DD/Y... h:mm:ss AM/PM 09/11/2023 12:00:00 AM

(UTC+05:30) Chennai, Kolkata... (UTC+05:30) Chennai, Kolkata...

OK Cancel

6. Open the URL (copy from Notepad) in browser and note that you are not able to download the file.

Summary

SAS Token

Account SAS

Always Ad-hoc

Same SAS token can be used for all services.

Service SAS

Ad-hoc / Policy based

Specific to only one service at a time.

Container SAS token works for all blobs in that container and can't be generated at portal

Blob SAS token is only for the blob for which it is generated

Azure Storage encryption for data at rest

- Azure Storage automatically encrypts your data when persisting it to the cloud by Storage Service Encryption (SSE)
- SSE automatically encrypts data when writing it to Azure Storage. When you read data from Azure Storage, Azure Storage decrypts the data before returning it.
- Storage accounts are encrypted regardless of their performance tier (standard or premium)
- All Azure Storage resources are encrypted, including blobs, disks, files, queues, and tables. All object metadata is also encrypted.
- Encryption does not affect Azure Storage performance.
- You can rely on **Microsoft-managed keys** for the encryption of your storage account, or you can manage encryption with your own keys.

Encryption in transit

- You can enable Encryption by setting **Secure transfer required**(Settings→Configuration)
- When enabled you need to use https for Rest API call and SMB 3.0 with encryption for Azure Files

Public Endpoint ,Service Endpoint and Private EndPoint

What is Public Endpoint?

For Many services in Azure, we have public endpoint, using that you can communicate with the service.

It is an address(Fully qualified domain name which can be communicated through internet. It is internet routable.

Azure offers two similar but distinct services to allow virtual network (VNet) resources to privately connect to other Azure services. [Azure VNet Service Endpoints](#) and [Azure Private Endpoints](#) .

Each network security by allowing VNet traffic to communicate with service resources without going over the internet

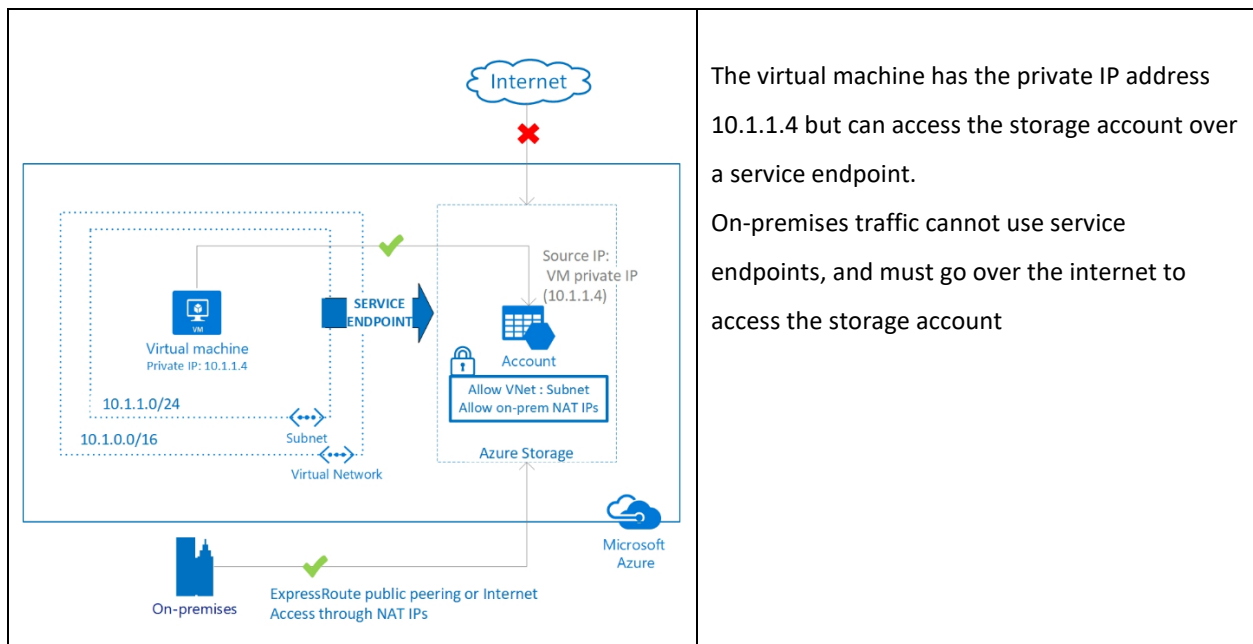
What is Service Endpoint?

- A [service endpoint](#) allows VNet resources to use private IP addresses to connect to an Azure service's public endpoint, meaning traffic flows to the service resource over the Azure backbone network - instead of over the internet.

- Service endpoints enable you to limit network access to Azure service resources via virtual network subnets and IP addresses you specify. Thus providing an extra layer of security allowing you to secure your critical Azure service resources to your virtual networks.
- You need to enable this on your subnet. It can be enabled for one or multiple services (Storage, Azure SQL etc).
- When it is enabled, Identity of virtual network will be available to services (Storage, Azure SQL etc).
- Service Endpoint provides security by specifying firewall at service level so that it only accepts traffic from the subnet associated with the service endpoint.

Set up service End Point

1. Enable service endpoint in subnet
2. Configure firewall rule at service level for that subnet in given virtual machine



Lab 10: Restrict access to Storage Account BLOB only from a VM in a given Subnet.

Prerequisites

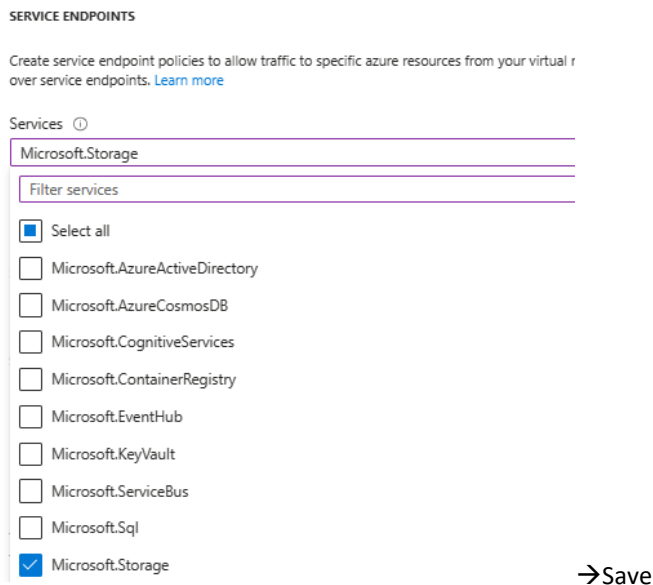
Create a virtual network

Create windows vm in it with rdp-allow and web-allow rule

1. Copy the URL of the uploaded file in storage account in public container(eg: <https://dssdemoaugsa.blob.core.windows.net/public-images/Tree.jpeg>)
2. Note that the URL is currently accessible from everywhere (from VM inside vNET and from my local machine outside vNET)

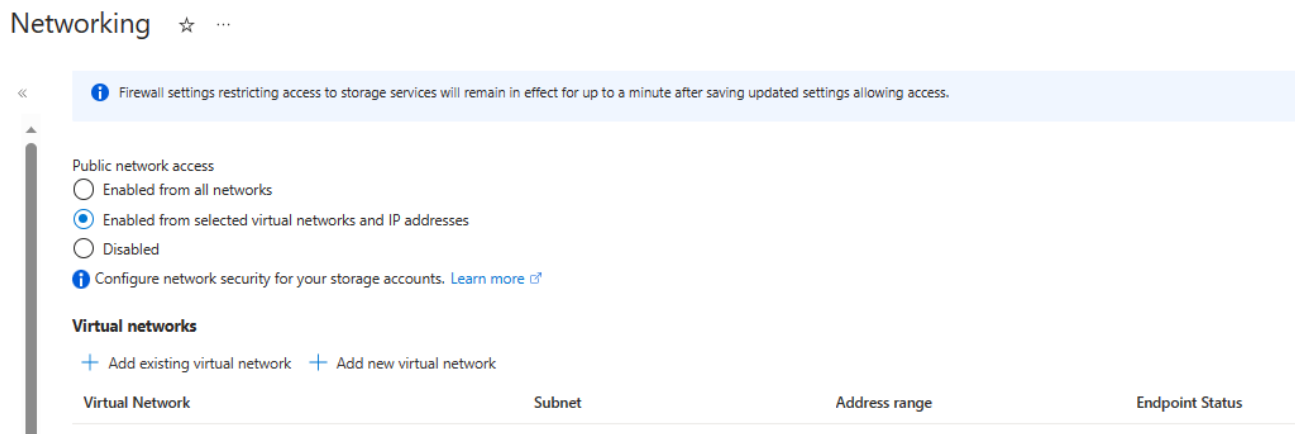
3. Enable Service Endpoint

DemoVnet → Settings → Subnet → Select(FrontEnd subnet) → Service endpoints → Select



4. Restrict access to Azure PaaS service (Storage Account) only from **FrontEnd** subnet.

- Storage Account → Networking → Firewalls and virtual networks Tab
- Public network access: Enabled from selected virtual networks and IP addresses
- Virtual networks → + Add existing virtual network



d. Save

2. Confirm access is denied to a resource from another subnet of vNET and also from the internet (my machine)

Your Local Machine → Browser → URL of BLOB uploaded to Storage Account

- Note that you **are not able** to view the blob/file content

3. Confirm access to a resource from an allowed subnet.

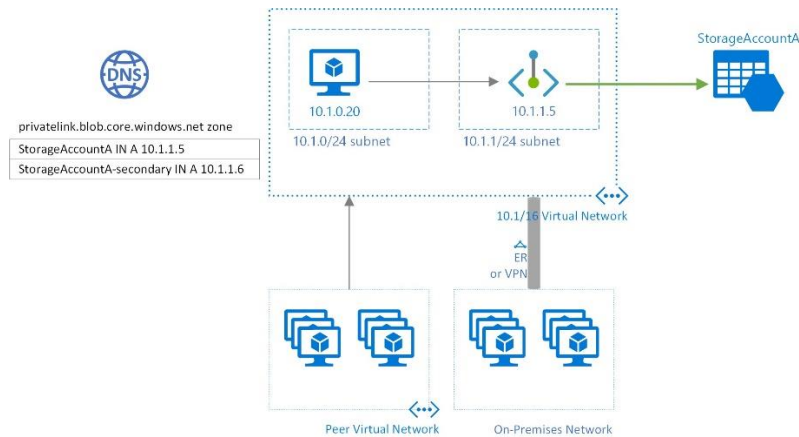
VM created in FrontEnd subnet → Browser → URL of BLOB uploaded to Storage Account

- a. Note that you are **able** to view the blob/file content.
4. VM→Launch command prompt
ping <storageaccount>.blob.core.windows.net
5. Observe the public ip for storage services.

Note: You can add your machine IP Address under firewall and access the Storage Account from your local machine.

What is Private End Point?

- Private End Point is a network interface that provides a private IP address to a service, that would normally only be accessible to a VNet via public IP address.
- Service resource (Storage etc) is extended in subnet in virtual network. The Resource is called Private Link Resource.
- NIC is created for that service and linked to service. Pass Service(Storage Account) also get private IP Address through NIC.
- Private endpoints enable traffic from on-premises to access private link resources without public peering or traversing the internet.
- So from On-Prem you can connect to Private link resource (Ex: Storage account) privately through Express Route OR site to site VPN
- When you create a private endpoint, the DNS CNAME resource record for the storage account is updated to an alias in a subdomain with the prefix privatelink.



Lab11: Ensure that Azure PaaS service uses Private IP, if requested from vNET and public IP, if requested from outside the vNET.

Prerequisites:

Create a Virtual Machine to test the private endpoint is working from it.

Create a Storage Account, Add Public Container and upload a file.

1. Create a **Private Endpoint**:

1. **Storage Account**→**Networking**→**Private endpoint connections**→**+Private endpoint**
Create a private endpoint ...

1 Basics 2 Resource 3 Virtual Network 4 DNS 5 Tags 6 Review + create

Use private endpoints to privately connect to a service or resource. Your private endpoint must be in the same region as your virtual network, but can be in a different region from the private link resource that you are connecting to. [Learn more](#)

Project details

Subscription * ⓘ Azure Training - SS1

Resource group * ⓘ DP-203-RG [Create new](#)

Instance details

Name * pvt-blob ✓

Network Interface Name * pvt-blob-nic ✓

Region * East US

2. →Next

3. **Target sub-resource = blob**

✓ Basics ✓ **Resource** 3 Virtual Network 4 DNS 5 Tags

Private Link offers options to create private endpoints for different Azure resources in an Azure storage account. Select which resource you would like to connect to use.

Subscription Azure Training - SS1 (24784a25-4b3b)

Resource type Microsoft.Storage/storageAccounts

Resource dssdemoaugsa

Target sub-resource * ⓘ blob

→Next

4. Select Virtual Network and Subnet in which it Private IP will be created.

- ✓ Basics ✓ Resource **3 Virtual Network** 4 DNS 5 Tags 6 Review + create

Networking

To deploy the private endpoint, select a virtual network subnet. [Learn more](#)

Virtual network *	<input type="text" value="DemoVnet (DP-203-RG)"/>
Subnet *	<input type="text" value="FrontEnd"/>
Network policy for private endpoints	Disabled (edit)

Private IP configuration


- ☒ Dynamically allocate IP address
☐ Statically allocate IP address

5. Next (DNS)→Next→ Review + Create
 6. This creates a Private DNS (<storageaccount>.blob.core.windows.net) name and NIC with Private IP address (192.168.0.5).
2. **Test connectivity to private endpoint from outside the VNet (your local machine)**
 1. nslookup <storageaccount>.blob.core.windows.net
 2. Note that the IP address listed is **public IP**
 3. Try Access to Blob storage file from local machine and note that is accessible*.
 3. ***To restrict public access to storage account**
 1. Go to Storage Account →Networking→ Firewalls and virtual networks → Select **Selected Networks** and don't select any network **OR Select Disabled**
 2. Try Access to Blob storage file from local machine and note that it is not accessible*.
 4. **Test connectivity to private endpoint from inside the VNet**
 1. RDP to VM
 2. nslookup <storageaccount>.blob.core.windows.net
 3. Note that the IP address listed is **private IP (It's the IP Address of the NIC created for Private endpoint)**
 4. Try Access to Blob storage file from VM and **it succeeds to connect.**

Azure File Storage

- Azure File storage is a service that offers file shares in the cloud using the standard [Server Message Block \(SMB\) Protocol](#).
- With Azure File storage, you can migrate **legacy applications** that rely on file shares to Azure quickly and without costly rewrites.

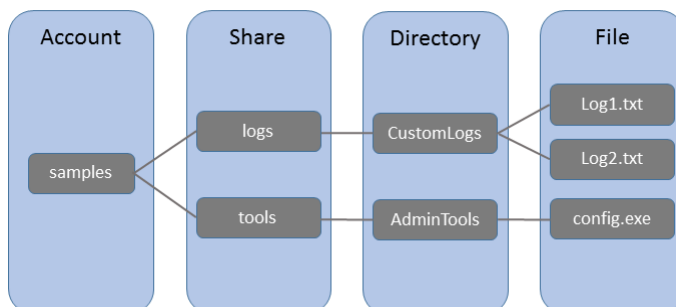
- Microsoft Azure virtual machines can share file data across application components via mounted shares, and on-premises applications can access file data in a share via the File storage API.
- You can control access to shares in Azure File Storage using authentication and authorization services available through Azure Active Directory Domain Services.
- It enables you to share up to 100 TB of data in a single storage account Which can be distributed across any number of file shares in the account.



Common uses of File storage include:

- Replace or supplement on-premises file servers.
- "Lift and shift" of Legacy applications, that rely on file shares to run on Azure virtual machines or cloud services, without expensive rewrites.
- Simplify cloud development.
 - Shared application settings, for example in configuration files.
 - Diagnostic share such as logs, metrics, and crash dumps in a shared location
 - Dev/Test/Debug
- Storing tools and utilities needed for developing or administering Azure virtual machines or cloud services

File storage concepts:



- **Storage Account:** All access to Azure Storage is done through a storage account.

- **Share:** A File storage share is an SMB file share in Azure. All directories and files must be created in a parent share. An account can contain an unlimited number of shares, and a share can store an unlimited number of files, up to the 5 TB total capacity of the file share.
- **Directory:** An optional hierarchy of directories.
- **File:** A file in the share. A file may be up to 1 TB in size but you can set quotas to limit the size of each share below this figure.
- **Max size of a File Share = 5TB**

URL format:

<https://<storage account>.file.core.windows.net/<share>/<directory/directory>/<file>>

<http://samples.file.core.windows.net/logs/CustomLogs/Log1.txt>

Lab12: Create File Share and Map file share in windows

Managing Using Azure Portal:

1. Storage Account → File shares → **+File share**
 - a. Name = myfiles
 - b. Review + create
 - c. Backup tab → Enable backup = Uncheck
 - d. Review + create → Create
2. Optionally add directory and upload the file[s].
 - a. myfiles → Upload
3. Connect to Fileshare from local machine
 - a. myfiles → **Connect**
 - i. Drive letter = Z
 - ii. Authentication mode = Storage account key
 - iii. Show script → Copy the script
4. Connect to Fileshare from local machine
 - a. Windows search → Powershell → run the script copied in previous step

```
$connectTestResult = Test-NetConnection -ComputerName dsdemostorage.file.core.windows.net -Port 445

if ($connectTestResult.TcpTestSucceeded) {
    # Save the password so the drive will persist on reboot
    cmd.exe /C "cmdkey /add:"dsdemostorage.file.core.windows.net" /user:"localhost\dsdemostorage"
    /pass:"DTLbeGfXlukUZt6svlb8dju1FW9kAn0XNLbshFKB4MX62FqqNzsDmD0S3bXofEsa1t8SxeCrO5q3+Ast
    hc0vXw==`""
    # Mount the drive
```



```
New-PSDrive -Name Z -PSProvider FileSystem -Root "\\dsdemostorage.file.core.windows.net\myfiles"
Persist
} else {
    Write-Error -Message "Unable to reach the Azure storage account via port 445. Check to make sure
your organization or ISP is not blocking port 445, or use Azure P2S VPN, Azure S2S VPN, or Express Route
to tunnel SMB traffic over a different port."
}
```

- b. Open to Z: in Windows Explorer and copy files into this folder
- c. Storage account → Files → myfiles → Browse → Note that the files are actually uploaded to file share.

Note: Ensure port 445 is open: Azure Files uses SMB protocol. SMB communicates over TCP port 445

Azcopy:

You can also use the *AzCopy copy* command to transfer files and folders to and from Azure File Storage to your local computer and between storage accounts.

Azure Table Storage

- The Azure Table storage service is not traditional table but Key/attribute(value) store.
- The Azure Table storage service stores large amounts of **partially structured** data offering high availability and massively scalable storage.
- The service is a **NoSQL datastore** which accepts **authenticated calls** from inside and outside the Azure cloud.

You can use the Table service to store and query huge sets of structured, non-relational data, and your tables will scale as demand increases.

- **Table:** A table is a collection of entities. Tables don't enforce a schema on entities, which means a single table can contain entities that have different sets of properties. The number of tables that a storage account can contain is limited only by the storage account capacity limit.
- **Entity:** An entity is a set of properties, similar to a database row. An entity can be up to **1MB in size**. In an Azure Table Storage table, items are referred to as *rows*, and fields are known as *columns*.
- **Properties:** A property is a name-value pair. Each entity can include up to **252 custom properties** to store up to **1 MB** of data. Each entity also has **3 system** properties that specify a **partition key** (string upto 1KB in size), a **row key** (string upto 1KB in size), and a **timestamp**. Entities with the same partition key can be **queried more quickly**, and inserted/updated in atomic operations. An entity's row key is its unique identifier within a partition.

Common uses of the Table service include:

- You can use Table storage to store flexible datasets, such as ,Product catalogues for eCommerce applications ,user data for web applications, address books, device information, and any other type of metadata that your service requires.
- Storing datasets that **don't** require complex joins, foreign keys, or stored procedures and can be de-normalized for fast access.
- Quickly querying data using a clustered index (Combination of PartitionKey and RowKey).

The URI for a specific table access is structured as follows:

<http://<account>.table.core.windows.net/<TableName>>

Lab13:Create table and Insert data.

Using Visual Studio Server Explorer

1. Server Explorer → ... →Select Windows Azure Storage → Select and Expand Storage Account.
2. Expand to Tables → Create Table..., Enter Name of Table
3. Select Table → Right click → View Table
4. Use the Editor to manage Table.